

# Basics

June 20, 2025

```
[1]: import pandas as pd

[2]: chunk_iter = pd.read_csv('crunchbase-investments.csv', chunksize=5000,
    ↪encoding='Latin-1')
    #we must specify the different encoding other than utf, it doesnt run as it
    ↪cant decode all the data.
```

0.0.1 Here the data is not just 5000, so we are going to iterate the chunks and then run the statement.

```
[3]: cb_length = []

    # looking at dataset's length
    for chunk in chunk_iter:
        cb_length.append(len(chunk))

    cb_length = sum(cb_length)

    print(f"Total number of rows in the dataset: {cb_length}")
```

Total number of rows in the dataset: 52870

0.0.2 Now we are going to re initialize and look for missing values.

```
[4]: # Counting missing values by column
    chunk_iter = pd.read_csv('crunchbase-investments.csv', chunksize=5000,
    ↪encoding='Latin-1')
    #empty list
    mv_list = []
    for chunk in chunk_iter:
        if not mv_list: # check if mv_list is empty
            columns = chunk.columns
            print(columns.sort_values())
            mv_list.append(chunk.isnull().sum())
    combined_mv_vc = pd.concat(mv_list)
    unique_combined_mv_vc = combined_mv_vc.groupby(combined_mv_vc.index).sum()
    unique_combined_mv_vc
```

```
Index(['company_category_code', 'company_city', 'company_country_code',
      'company_name', 'company_permalink', 'company_region',
      'company_state_code', 'funded_at', 'funded_month', 'funded_quarter',
      'funded_year', 'funding_round_type', 'investor_category_code',
      'investor_city', 'investor_country_code', 'investor_name',
      'investor_permalink', 'investor_region', 'investor_state_code',
      'raised_amount_usd'],
      dtype='object')
```

```
[4]: company_category_code      643
      company_city              533
      company_country_code       1
      company_name               1
      company_permalink           1
      company_region             1
      company_state_code        492
      funded_at                  3
      funded_month               3
      funded_quarter             3
      funded_year                3
      funding_round_type         3
      investor_category_code    50427
      investor_city             12480
      investor_country_code     12001
      investor_name              2
      investor_permalink         2
      investor_region            2
      investor_state_code       16809
      raised_amount_usd         3599
      dtype: int64
```

As we can see, we have around 52k rows and the investor category code has around 50k empty columns which makes it a good choice to drop the column.

Now lets look at the memory usage of the columns.

```
[5]: # Total memory footprint for each column

chunk_iter = pd.read_csv('crunchbase-investments.csv', chunksize=5000,
    encoding='Latin-1')
counter = 0
series_memory_fp = pd.Series(dtype='float64')
for chunk in chunk_iter:
    if counter == 0:
        series_memory_fp = chunk.memory_usage(deep=True)
    else:
        series_memory_fp += chunk.memory_usage(deep=True)
```

```
counter += 1
```

```
# Drop memory footprint calculation for the index.
```

```
series_memory_fp_before = series_memory_fp.drop('Index').sort_values()
```

```
series_memory_fp_before
```

```
[5]: funded_year          422960
      raised_amount_usd    422960
      investor_category_code 622424
      investor_state_code   2476607
      investor_country_code 2647292
      investor_city         2885083
      company_state_code    3106051
      company_country_code  3172176
      funded_month         3383584
      funded_quarter        3383584
      investor_region       3396281
      funding_round_type    3410707
      company_region        3411545
      company_category_code 3421104
      company_city          3505886
      funded_at             3542185
      company_name          3591326
      investor_name         3915666
      company_permalink     4057788
      investor_permalink    4980548
      dtype: int64
```

checking the total memory usage!

```
[6]: # Total memory footprint of the data (in megabytes)
      series_memory_fp_before.sum() / (1024 * 1024)
```

```
[6]: np.float64(56.98753070831299)
```

```
[7]: # Drop columns representing URLs or containing too many missing values (>90%
      ↪missing)
      drop_cols = ['investor_permalink', 'company_permalink',
      ↪'investor_category_code', 'funded_month', 'funded_quarter', 'funded_year']
      keep_cols = chunk.columns.drop(drop_cols)
```

```
[8]: keep_cols.tolist
```

```
[8]: <bound method IndexOpsMixin.tolist of Index(['company_name',
      'company_category_code', 'company_country_code',
      'company_state_code', 'company_region', 'company_city', 'investor_name',
      'investor_country_code', 'investor_state_code', 'investor_region',
```

```

        'investor_city', 'funding_round_type', 'funded_at',
        'raised_amount_usd'],
        dtype='object')>

```

### 0.0.3 Selecting the data types

```

[9]: # Key: Column name, Value: List of types
col_types = {}
chunk_iter = pd.read_csv('crunchbase-investments.csv', chunksize=5000,
    encoding='Latin-1', usecols=keep_cols)

for chunk in chunk_iter:
    for col in chunk.columns:
        if col not in col_types:
            col_types[col] = [str(chunk.dtypes[col])]
        else:
            col_types[col].append(str(chunk.dtypes[col]))

```

```

[10]: uniq_col_types = {}
for k,v in col_types.items():
    uniq_col_types[k] = set(col_types[k])
uniq_col_types

```

```

[10]: {'company_name': {'object'},
      'company_category_code': {'object'},
      'company_country_code': {'object'},
      'company_state_code': {'object'},
      'company_region': {'object'},
      'company_city': {'object'},
      'investor_name': {'object'},
      'investor_country_code': {'float64', 'object'},
      'investor_state_code': {'float64', 'object'},
      'investor_region': {'object'},
      'investor_city': {'float64', 'object'},
      'funding_round_type': {'object'},
      'funded_at': {'object'},
      'raised_amount_usd': {'float64'}}

```

```

[11]: chunk.head(3)

```

```

[11]:
      company_name company_category_code company_country_code \
50000      NuORDER              fashion              USA
50001      ChaCha              advertising              USA
50002      Binfire              software              USA

      company_state_code company_region  company_city  investor_name \
50000              CA      Los Angeles  West Hollywood  Mortimer Singer

```

50001	IN	Indianapolis	Carmel	Morton Meyerson
50002	FL	Bocat Raton	Bocat Raton	Moshe Ariel

	investor_country_code	investor_state_code	investor_region	\
50000	NaN	NaN	unknown	
50001	NaN	NaN	unknown	
50002	NaN	NaN	unknown	

	investor_city	funding_round_type	funded_at	raised_amount_usd
50000	NaN	series-a	2012-10-01	3060000.0
50001	NaN	series-b	2007-10-01	12000000.0
50002	NaN	angel	2008-04-18	500000.0

```
[13]: # Analyzing unique value counts
unique_values = {}

chunk_iter = pd.read_csv('crunchbase-investments.csv', chunksize=5000,
    encoding='Latin-1', usecols=keep_cols)
for chunk in chunk_iter:
    for col in chunk.columns:
        if col not in unique_values:
            unique_values[col] = set()
            unique_values[col].update(chunk[col].unique())

for col, unique_vals in unique_values.items():
    print(f"Column: {col}, Unique Values: {len(unique_vals)}")
```

```
Column: company_name, Unique Values: 11574
Column: company_category_code, Unique Values: 44
Column: company_country_code, Unique Values: 3
Column: company_state_code, Unique Values: 51
Column: company_region, Unique Values: 547
Column: company_city, Unique Values: 1230
Column: investor_name, Unique Values: 10466
Column: investor_country_code, Unique Values: 75
Column: investor_state_code, Unique Values: 53
Column: investor_region, Unique Values: 586
Column: investor_city, Unique Values: 993
Column: funding_round_type, Unique Values: 10
Column: funded_at, Unique Values: 2809
Column: raised_amount_usd, Unique Values: 1469
```

```
[14]: col_types = {
    'company_category_code': 'category',
    'funding_round_type': 'category',
    'investor_state_code': 'category',
    'investor_country_code': 'category'
```

```

}

chunk_iter = pd.read_csv('crunchbase-investments.csv',
                          chunksize=5000,
                          encoding='Latin-1',
                          usecols=keep_cols,
                          dtype=col_types,
                          parse_dates=['funded_at'])

counter = 0
series_memory_fp = pd.Series(dtype='float64')
for chunk in chunk_iter:
    if counter == 0:
        series_memory_fp = chunk.memory_usage(deep=True)
    else:
        series_memory_fp += chunk.memory_usage(deep=True)
    counter += 1

# Drop memory footprint calculation for the index.
series_memory_fp_after = series_memory_fp.drop('Index').sort_values()
print(series_memory_fp_after)

# Total memory footprint of the data (in megabytes)
series_memory_fp_after.sum() / (1024 * 1024)

```

```

funding_round_type      61906
investor_country_code   81118
investor_state_code     81811
company_category_code   94160
funded_at               422960
raised_amount_usd       422960
investor_city           2885083
company_state_code      3106051
company_country_code    3172176
investor_region         3396281
company_region          3411545
company_city            3505886
company_name            3591326
investor_name           3915666
dtype: int64

```

```
[14]: np.float64(26.844910621643066)
```

```

[15]: df = pd.DataFrame({
        'Before': series_memory_fp_before,
        'After': series_memory_fp_after,
    })

```

```

    'Difference': series_memory_fp_after - series_memory_fp_before
})

df.sort_values('Difference')

```

```

[15]:

```

	Before	After	Difference
funding_round_type	3410707	61906.0	-3348801.0
company_category_code	3421104	94160.0	-3326944.0
funded_at	3542185	422960.0	-3119225.0
investor_country_code	2647292	81118.0	-2566174.0
investor_state_code	2476607	81811.0	-2394796.0
company_country_code	3172176	3172176.0	0.0
company_name	3591326	3591326.0	0.0
company_city	3505886	3505886.0	0.0
company_region	3411545	3411545.0	0.0
company_state_code	3106051	3106051.0	0.0
investor_name	3915666	3915666.0	0.0
investor_city	2885083	2885083.0	0.0
investor_region	3396281	3396281.0	0.0
raised_amount_usd	422960	422960.0	0.0
company_permalink	4057788	NaN	NaN
funded_month	3383584	NaN	NaN
funded_quarter	3383584	NaN	NaN
funded_year	422960	NaN	NaN
investor_category_code	622424	NaN	NaN
investor_permalink	4980548	NaN	NaN

```

[16]: print(f'Before total memory: {series_memory_fp_before.sum() / (1024 * 1024)}')
      print(f'After total memory: {series_memory_fp_after.sum() / (1024 * 1024)}')

```

Before total memory: 56.98753070831299  
 After total memory: 26.844910621643066

#### 0.0.4 Loading Chunks into SQLite

```

[19]: #import SQLite3
      import sqlite3
      #this creates a db in jyp env
      conn = sqlite3.connect('crunchbase.db')
      cursor = conn.cursor()

```

```

[20]: chunk_iter = pd.read_csv('crunchbase-investments.csv',
                               chunksize=5000,
                               encoding='Latin-1',
                               usecols=keep_cols,
                               dtype=col_types,
                               parse_dates=['funded_at'])

```

```

    )
for chunk in chunk_iter:
    chunk.to_sql("investments", conn, if_exists='append', index=False)

cursor.execute("SELECT name FROM sqlite_master WHERE type='table';")
tables = cursor.fetchall()

if tables:
    print("Tables found in the database:", [t[0] for t in tables])
else:
    print("No tables found in the database.")

```

Tables found in the database: ['investments']

```
[21]: cursor.execute("SELECT * FROM investments LIMIT 5;")
print(cursor.fetchall())
```

```

[('AdverCar', 'advertising', 'USA', 'CA', 'SF Bay', 'San Francisco',
'1-800-FLOWERS.COM', 'USA', 'NY', 'New York', 'New York', 'series-a',
'2012-10-30 00:00:00', 2000000.0), ('LaunchGram', 'news', 'USA', 'CA', 'SF Bay',
'Mountain View', '10Xelerator', 'USA', 'OH', 'Columbus', 'Columbus', 'other',
'2012-01-23 00:00:00', 20000.0), ('uTaP', 'messaging', 'USA', None, 'United
States - Other', None, '10Xelerator', 'USA', 'OH', 'Columbus', 'Columbus',
'other', '2012-01-01 00:00:00', 20000.0), ('ZoopShop', 'software', 'USA', 'OH',
'Columbus', 'columbus', '10Xelerator', 'USA', 'OH', 'Columbus', 'Columbus',
'angel', '2012-02-15 00:00:00', 20000.0), ('eFuneral', 'web', 'USA', 'OH',
'Cleveland', 'Cleveland', '10Xelerator', 'USA', 'OH', 'Columbus', 'Columbus',
'other', '2011-09-08 00:00:00', 20000.0)]

```

### 0.0.5 A gist with visualization of our data

```
[23]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

pd.set_option('display.float_format', '{:,.2f}'.format)

# Query and prepare data
df = pd.read_sql("""
    SELECT company_category_code, AVG(raised_amount_usd)
    FROM investments
    GROUP BY company_category_code
    ORDER BY AVG(raised_amount_usd) DESC;
    """, conn)

df = df.set_index('company_category_code')
```



```

# Select top N categories
top_n = 15
top_df = df.head(top_n)

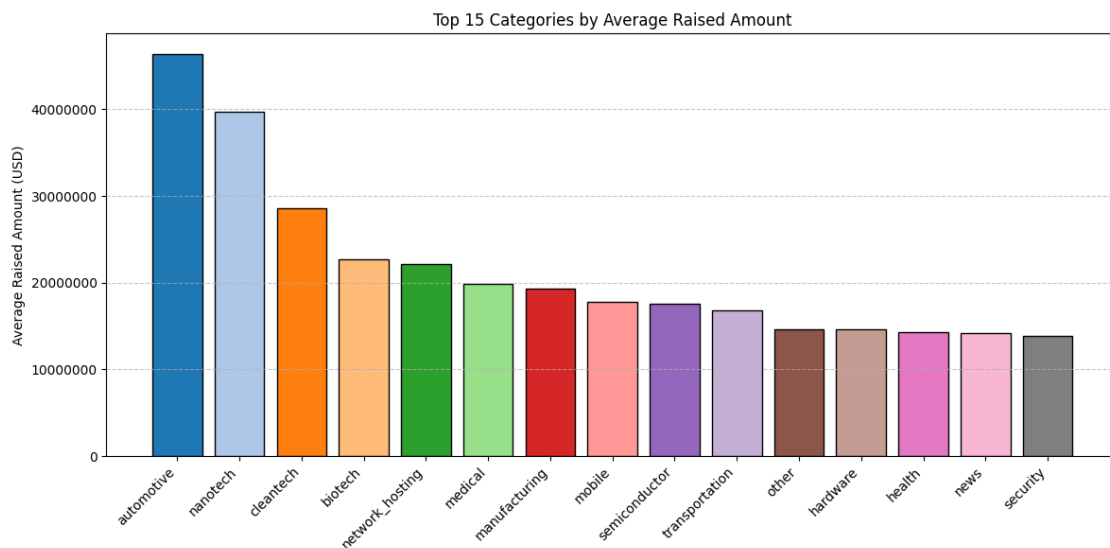
# Generate distinct colors for each bar
colors = sns.color_palette('tab20', n_colors=top_n)

# Plot with unique colors for each bar
fig, ax = plt.subplots(figsize=(12, 6))
bars = ax.bar(top_df.index, top_df['AVG(raised_amount_usd)'], color=colors,
               edgecolor='black')

# Customize the plot
ax.set_ylabel('Average Raised Amount (USD)')
ax.set_title(f'Top {top_n} Categories by Average Raised Amount')
ax.ticklabel_format(style='plain', axis='y')
plt.xticks(rotation=45, ha='right')
ax.grid(True, axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()

```



[ ]: